

Golang Estonia

**Production
Ready
Concurrency**



Golang Estonia Meetup Group

By the community, for the community

- Build a community
- Give professional content
- Share experiences
- Foster professional development
- Support interaction with experts

Get Involved

- *Join Gopher Slack #estonia*
- *Become a Speaker*
 - *Lightning talks (5-10 min)*
 - *Workshops*
 - *Talks (~20 min)*
- *Host (talk to your company)*
- *Become a Sponsor (e.g. sponsor mugs, raffles)*
- *Become a Co-Organizer*

Production Ready Concurrency

Egon Elbre
@egonelbre

Disclaimer

The following describes how to write a system where it's hard to make mistakes and easy to test.

This occasionally comes at the cost of performance and lines of code.

*PS: There are plenty of exceptions,
I will skip over most of them.*

Outline

Rules of thumb

Choosing primitives

A few custom primitives

Q & A

Rules of Thumb

No global variables

including global flags, caches, pools, logging etc. make them scoped
they all complicate testing

Use concurrency only when needed




```
var wg sync.WaitGroup
```

```
wg.Add(1)
```

```
go serve()
```

```
wg.Wait()
```



```
serve()
```

System without concurrency is much easier to debug, test and understand.

but, all your tests should use `t.Parallel()`




```
func TestXYZ(t *testing.T) {  
    t.Parallel()  
    ...  
}
```

```
go test -race ./...
```

Prefer synchronous API



```
server.Start(ctx)  
server.Stop()  
server.Wait()
```




```
server.Run(ctx)
```

It's obvious how to turn sync to async, but the reverse isn't always true.

Know when things stop



```
go listenHTTP()  
go listenGRPC()  
go listenDebug()  
select {}
```




```
var endpoints errgroup.Group  
endpoints.Go(listenHTTP)  
endpoints.Go(listenGRPC)  
endpoints.Go(listenDebug)  
return endpoints.Wait()
```

At the end of your tests, all goroutines you started should have finished.

Cancellation-Aware



```
time.Sleep(time.Second)
```

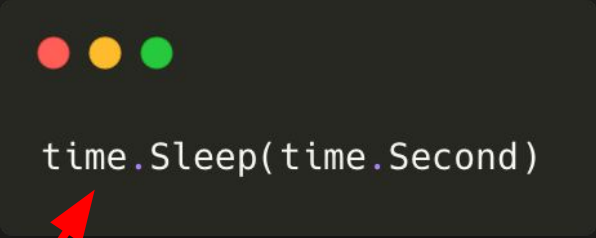


```
tick := time.NewTimer(time.Second)
defer tick.Stop()

select {
case <-tick.C:
case <-ctx.Done():
    return ctx.Err()
}
```

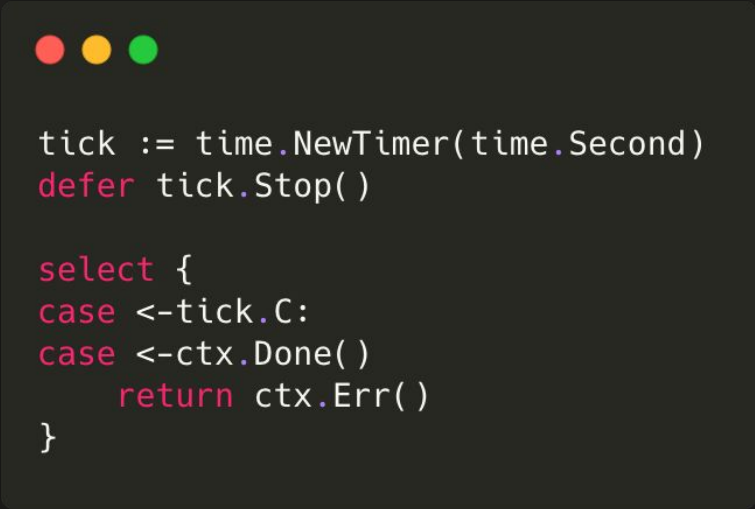

Cancellation handling is required to have everything shutdown fast on Ctrl-C

Cancellation-Aware



```
time.Sleep(time.Second)
```

often indicates a problem



```
tick := time.NewTimer(time.Second)
defer tick.Stop()


select {
case <-tick.C:
case <-ctx.Done():
    return ctx.Err()
}
```

Context handling is required to have everything shutdown fast on Ctrl-C

Cancellation-Aware - Part 2



```
for _, f := range files {  
    data, err := os.ReadFile(f)  
    ...  
}
```




```
for _, f := range files {  
    if err := ctx.Err(); err != nil {  
        return err  
    }  
    data, err := os.ReadFile(f)  
    ...  
}
```

Anything taking more than 50ms should be ctx.Err aware.

Avoid worker pools

```
work := make(chan string, 10)
for k := 0; k < 10; k++ {
    go func () {
        for _, w := range work {
            process(w)
        }
    }()
}

for _, item := range items {
    work <- item
}
```




```
limiter := newLimiter(10)
for _, item := range items {
    item := item
    limiter.Go(func() {
        process(item)
    })
}
limiter.Wait()
```

Worker pools make stack traces harder to read, use more resources and are often slower.

Avoid polling state



```
lastKnown := 0
for {
    time.Sleep(time.Second)
    t.mu.Lock()
    if lastKnown != t.current {
        process(t.current)
        lastKnown = t.current
    }
    t.mu.Unlock()
}
```



This would be a whole talk

Wastes resources and makes go runtime work harder.
Also it's slower to respond to changes to the value.

Defer post-actions (unlock, wait, close)



```
for _, item := range items {  
    service.mu.Lock()  
    service.process(item)  
    service.mu.Unlock()  
}
```




```
for _, item := range items {  
    func() {  
        service.mu.Lock()  
        defer service.mu.Unlock()  
        service.process(item)  
    }()  
}
```

It's clearer when a Unlock is missing.
Code modifications are less likely to introduce a bug.

Don't expose your locks



```
type Set struct {  
    sync.Lock  
    ...  
}
```



```
type Set struct {  
    mu sync.Lock  
    ...  
}
```

This is a quick way to get deadlocks.

Name your goroutines



```
labels := pprof.Labels("worker", "purge")
pprof.Do(ctx, labels,
    func(ctx context.Context) {
        ...
    })
```

Significantly helps with debugging.
See more at <https://rakyll.org/profiler-labels/>

Rules of Thumb: Summary

Use only when required

Prefer Synchronous API (async can be optional)

Know when things stop

Context aware code

No worker pools & polling

Defer post-actions

Don't expose your locks

Name your goroutines

Choosing your

concurrency primitives

Order of preference

1. no-concurrency
2. golang.org/x/sync, `sync.Once`
3. custom solution (or library)
4. `sync.Mutex` (special cases)

↓↓↓ *only for implementing custom solutions* ↓↓↓

5. `sync.Map`, `sync.Pool` (need a typesafe wrapper)
6. `sync.WaitGroup`
7. `chan`, `go func() {`
8. `sync.Mutex`, `sync.Cond`
9. `atomics`

Problem #1: go func() {



```
func (server *Server) ServeHTTP(w http.ResponseWriter, r *http.Request) {  
    ...  
    go func() {  
        res, err := server.db.ExecContext(r.Context(), "INSERT ...")  
        ...  
    }()  
    ...  
}  
  
func main() {  
    db, err := openDB()  
    defer db.Close()  
  
    err := server.Run(ctx)  
    ...  
}
```


Problem #1: go func() {

```
func (server *Server) ServeHTTP(w http.ResponseWriter, r *http.Request) {  
    ...  
    go func() {  
        res, err := server.db.ExecContext(r.Context(), "INSERT ...")  
        ...  
    }()  
    ...  
}  
  
func main() {  
    db, err := openDB()  
    defer db.Close()  
  
    err := server.Run(ctx)  
    ...  
}
```



'db' might be closed here

Problem #1: go func() {

```
func (server *Server) ServeHTTP(w http.ResponseWriter, r *http.Request) {  
    ...  
    go func() {  
        res, err := server.db.ExecContext(r.Context(), "INSERT ...")  
        ...  
    }()  
    ...  
}  
  
func main() {  
    db, err := openDB()  
    defer db.Close()  
  
    err := server.Run(ctx)  
    ...  
}
```




This context will get cancelled.

Problem #2: sync.WaitGroup



```
func processConcurrently(item []*Item) {  
    var wg sync.WaitGroup  
    for _, item := range items {  
        item := item  
        go func() { process(&wg, item) }()  
    }  
    wg.Wait()  
}  
  
func process(wg *sync.WaitGroup, item *Item) {  
    wg.Add(1)  
    defer wg.Done()  
  
    ...  
}
```

Problem #2: sync.WaitGroup



```
func processConcurrently(item []*Item) {  
    var wg sync.WaitGroup  
    for _, item := range items {  
        item := item  
        go func() { process(&wg, item) }()  
    }  
    wg.Wait()  
}
```

```
func process(wg *sync.WaitGroup, item *Item) {  
    wg.Add(1)  
    defer wg.Done()  
  
    ...  
}
```

logically racing



Problem #3: sync.WaitGroup




```
var wg sync.WaitGroup
wg.Add(len(items))

for _, item := range items {
    item := item
    if filepath.Ext(item.Path) != ".go" {
        continue
    }
    go func() {
        defer wg.Done()
        process(item)
    }()
}

wg.Wait()
```

Problem #3: sync.WaitGroup



```
var wg sync.WaitGroup
wg.Add(len(items))

for _, item := range items {
    item := item
    if filepath.Ext(item.Path) != ".go" {
        continue
    }
    go func() {
        defer wg.Done()
        process(item)
    }()
}

wg.Wait()
```

these can get out of sync



Tracking primitives



```
func (server *Server) ServeHTTP(w http.ResponseWriter, r *http.Request) {  
    ...  
    err := server.pending.Go(r.Context(), func(ctx context.Context) {  
        res, err := server.db.ExecContext(r.Context(), "INSERT ...")  
        ...  
    })  
    ...  
}  
  
func (server *Server) Run(ctx context.Context) error {  
    ...  
    defer server.pending.Wait()  
    ...  
}
```

Tracking primitives



```
func (server *Server) ServeHTTP(w http.ResponseWriter, r *http.Request) {  
    ...  
    err := server.pending.Go(r.Context(), func(ctx context.Context) {  
        ...  
        err := server.db.ExecContext(r.Context(), "INSERT ...")  
        ...  
    })  
    ...  
}
```

When server is shutting down
we should respond accordingly.

```
func (server *Server) Run(ctx context.Context) error {  
    ...  
    defer server.pending.Wait()  
    ...  
}
```


golang.org/x/sync/errgroup

waits all to stop on their own



```
var g errgroup.Group
g.Go(func() error {
    return public.Run(ctx)
})
g.Go(func() error {
    return private.Run(ctx)
})
err := g.Wait()
```

Cancels ctx when one fails



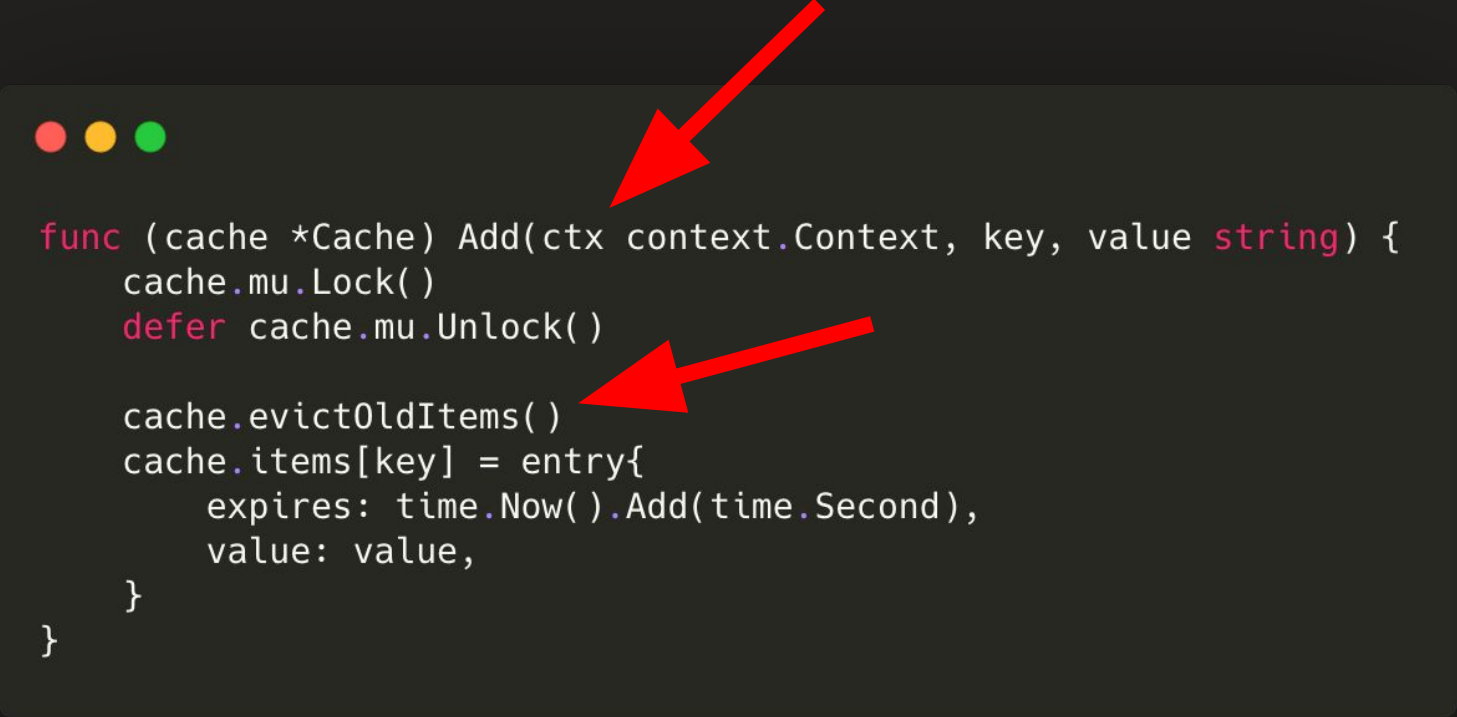
```
g, ctx := errgroup.WithContext(ctx)
g.Go(func() error {
    return public.Run(ctx)
})
g.Go(func() error {
    return private.Run(ctx)
})
err := g.Wait()
```

Problem #4: sync.Mutex



```
func (cache *Cache) Add(ctx context.Context, key, value string) {  
    cache.mu.Lock()  
    defer cache.mu.Unlock()  
  
    cache.evictOldItems()  
    cache.items[key] = entry{  
        expires: time.Now().Add(time.Second),  
        value: value,  
    }  
}
```

Problem #4: sync.Mutex



```
func (cache *Cache) Add(ctx context.Context, key, value string) {  
    cache.mu.Lock()  
    defer cache.mu.Unlock()  
  
    cache.evictOldItems()  
    cache.items[key] = entry{  
        expires: time.Now().Add(time.Second),  
        value: value,  
    }  
}
```

Alternative to sync.Mutex

make(chan *Item, 1)



```
func(cache *Cache) Add(ctx context.Context, key, value string) error {
    select {
    case <-ctx.Done():
        return ctx.Err()

    case state := <-cache.state:
        defer func() { cache.state <- state }()
        state.evictOldItems()

        state.items[key] = entry{
            expires: time.Now().Add(time.Second),
            value: value
        }
        return nil
    }
}
```

sync.Mutex

- only for protecting reading/writing to variables
- only in cases where code is $O(N)$ or faster and N is small

Main issues:

- doesn't allow responding to ctx cancellation
- many nested locks is a quick way to hit lock inversions

sync.RWMutex


(similar to sync.Mutex, but worse)

- only use when sync.Mutex is **demonstrably** slower
 - sync.Mutex is almost always faster

additional issues:

- *when you have lots of readers and no writers, there's still cache contention between the readers because taking a read lock mutates the mutex (not scalable)*
- *a writer attempting to grab the lock blocks future readers from acquiring it. so long lived readers with infrequent writers causes long delays of no work*

Problem #5: chan



```
const workerCount = 100

var wg sync.WaitGroup
workQueue := make(chan *Item)

for i := 0; i < workerCount; i++ {
    wg.Add(1)
    go func() {
        defer wg.Done()
        for item := range workQueue {
            process(item)
        }
    }()
}

err := db.IterateItems(ctx, func(item *Item) {
    workQueue <- item
})

wg.Wait()
```

Problem #5: chan

workQueue is not closed

```
const workerCount = 100

var wg sync.WaitGroup
workQueue := make(chan *Item)

for i := 0; i < workerCount; i++ {
    wg.Add(1)
    go func() {
        defer wg.Done()
        for item := range workQueue {
            process(item)
        }
    }()
}

err := db.IterateItems(ctx, func(item *Item) {
    workQueue <- item
})

wg.Wait()
```


Order of preference

1. no-concurrency
2. golang.org/x/sync, `sync.Once`
3. custom solution (or library)
4. `sync.Mutex` (special cases)

↓↓↓ *only for implementing custom solutions* ↓↓↓

5. `sync.Map`, `sync.Pool` (need a typesafe wrapper)
6. `sync.WaitGroup`
7. `chan`, `go func() {`
8. `sync.Mutex`, `sync.Cond`
9. `atomics`

My rule of thumb: use these only in concurrency primitives

- `make(chan X, N)`
- `go func() {`
- `sync.WaitGroup`

they are error-prone, hard to test, difficult to review

** but they are also not "the-end-of-the-world bad"*

note: using `select` is fine

A few custom primitives

There are many variations.

You need to figure out what is
best for your codebase.


Batch processing a slice

Batch processing a slice



```
func Parallel(n int, batchSize int, process func(low, high int)) {  
    var wg sync.WaitGroup  
    defer wg.Wait()  
  
    if batchSize <= 0 {  
        batchSize = 1  
    }  
  
    for low := 0; low < n; low += batchSize {  
        low, high := low, low + batchSize  
        if high > n {  
            high = n  
        }  
  
        wg.Add(1)  
        go func() {  
            defer wg.Done()  
            process(low, high)  
        }()  
    }  
}
```

Batch processing a slice



```
func Parallel(n int, batchSize int, process func(low, high int)) {
    var wg sync.WaitGroup
    defer wg.Wait()

    if batchSize <= 0 {
        batchSize = 1
    }

    for low := 0; low < n; low += batchSize {
        low, high := low, low + batchSize
        if high > n {
            high = n
        }

        wg.Add(1)
        go func() {
            defer wg.Done()
            process(low, high)
        }()
    }
}
```



```
var mu sync.Mutex
total := 0
```

```
Parallel(len(items), 256, func(low, high int) {
    price := 0
    for _, item := range items[low:high] {
        price += item.Price
    }

    mu.Lock()
    defer mu.Unlock()
    total += price
})
```

*this example is
probably not faster
than single core*

Sleeping

Sleeping




```
func Sleep(ctx context.Context, dur time.Duration) error {  
    t := time.NewTicker(dur)  
    defer t.Stop()  
  
    select {  
    case <-t.C:  
        return nil  
    case <-ctx.Done():  
        return ctx.Err()  
    }  
}
```

Sleeping



```
func Sleep(ctx context.Context, dur time.Duration) error {  
    t := time.NewTicker(dur)  
    defer t.Stop()  
  
    select {  
    case <-t.C:  
        return nil  
    case <-ctx.Done():  
        return ctx.Err()  
    }  
}
```



```
if err := Sleep(ctx, time.Second); err != nil {  
    return err  
}
```

Retrying with backoff

Retrying with backoff



```
retry := NewRetry(10, time.Second/10, time.Second)
for retry.Next(ctx) {

}
if err := retry.Err(); err != nil {
    return err
}
```

```
), time.Second/10,
.Context() error {
```

...

```
})
```

implementation is an exercise for the listener

Retrying with backoff



```
retry := NewRetry(10, time.Second/10, time.Second)
for retry.Next(ctx) {

}
if err := retry.Err(); err != nil {
    return err
}
```



```
err := Retry(ctx, 10, time.Second/10,
    func(ctx context.Context) error {
        ...
    })
```

implementation is an exercise for the listener

Running few things concurrently

Running few things concurrently



```
func Concurrently(
    ctx context.Context,
    fns ...func(context.Context) error,
) error {
    var g errgroup.Group
    for _, fn := range fns {
        fn := fn
        g.Go(func() error { return fn(ctx) })
    }
    return g.Wait()
}
```

Running few things concurrently



```
func Concurrently(
    ctx context.Context
    fns ...func(context.Context) error
) error {
    var g errgroup.Group
    for _, fn := range fns {
        g.Go(func() error {
            return fn(ctx)
        })
    }
    return g.Wait()
}
```



```
err := Concurrently(ctx,
    func(context.Context) error {
        // ...
    },
    func(context.Context) error {
        // ...
    },
    func(context.Context) error {
        // ...
    },
)
```


Running few things concurrently



```
func Concurrently(  
    ctx context.Context  
    fns ...func(context.Context) error  
) error {  
    var g errgroup.Group  
    for _, fn := range fns  
        fn := fn  
        g.Go(func() error {  
            return fn(ctx)  
        })  
    return g.Wait()  
}
```



```
err := Concurrently(ctx,  
    func(ctx context.Context) error {  
        // ...  
    },  
    func(ctx context.Context) error {  
        // ...  
    },  
    func(ctx context.Context) error {  
        // ...  
    },  
)
```



what should happen
when only one fails?

Running few things concurrently



```
func Concurrently(  
    ctx context.Context  
    fns ...func(context.Context) error  
) error {  
    var g errgroup.Group  
    for _, fn := range fns {  
        fn := fn  
        g.Go(func() error {  
            return fn(ctx)  
        })  
    }  
    return g.Wait()  
}
```



```
err := Concurrently(ctx,  
    func(ctx context.Context) error {  
        // ...  
    },  
    func(ctx context.Context) error {  
        // ...  
    },  
    func(ctx context.Context) error {  
        // ...  
    },  
)
```

what should happen
when only one fails?

what should happen
when one or more panics?

Waiting for a thing

Waiting for a thing



```
type Fence struct {  
    create sync.Once  
    release sync.Once  
    wait    chan struct{}  
}  
  
func (f *Fence) init() {  
    f.create.Do(func() {  
        f.wait = make(chan struct{})  
    })  
}
```



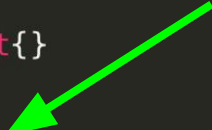
```
func (f *Fence) Release() {  
    f.init()  
    f.release.Do(func(){  
        close(f.wait)  
    })  
}  
  
func (f *Fence) Released() chan struct{} {  
    f.init()  
    return f.wait  
}  
  
func (f *Fence) Wait(ctx context.Context) error {  
    f.init()  
    select {  
    case <-f.Released():  
        return nil  
    case <-ctx.Done():  
        return ctx.Err()  
    }  
}
```

Waiting for a thing

```
type Fence struct {  
    create sync.Once  
    release sync.Once  
    wait    chan struct{}  
}
```

```
func (f *Fence) init() {  
    f.create.Do(func() {  
        f.wait = make(chan struct{})  
    })  
}
```

doesn't require
a constructor



```
func (f *Fence) Release() {  
    f.init()  
    f.release.Do(func(){  
        close(f.wait)  
    })  
}
```

```
func (f *Fence) Released() chan struct{} {  
    f.init()  
    return f.wait  
}
```

```
func (f *Fence) Wait(ctx context.Context) error {  
    f.init()  
    select {  
    case <-f.Released():  
        return nil  
    case <-ctx.Done():  
        return ctx.Err()  
    }  
}
```

Waiting for a thing



```
type Fence struct {  
    create sync.Once  
    release sync.Once  
    wait    chan struct{}  
}  
  
func (f *Fence) init() {  
    f.create.Do(func() {  
        f.wait = make(chan struct{})  
    })  
}
```



```
func (f *Fence) Release() {  
    f.init()  
    f.release.Do(func(){  
        close(f.wait)  
    })  
}
```

```
func (f *Fence) Released() chan struct{} {  
    f.init()  
    return f.wait  
}
```

```
func (f *Fence) Wait(ctx context.Context) error {  
    f.init()  
    select {  
    case <-f.Released():  
        return nil  
    case <-ctx.Done():  
        return ctx.Err()  
    }  
}
```

in some cases
you could want
a value to be passed
along

Waiting for a thing



```
type Fence struct {  
    create sync.Once  
    release sync.Once  
    wait    chan struct{}  
}  
  
func (f *Fence) init() {  
    f.create.Do(func() {  
        f.wait = make(chan struct{})  
    })  
}
```



```
func (f *Fence) Release() {  
    f.init()  
    f.release.Do(func(){  
        close(f.wait)  
    })  
}  
  
func (f *Fence) Released() chan struct{} {  
    f.init()  
    return f.wait  
}  
  
func (f *Fence) Wait(ctx context.Context) error {  
    f.init()  
    select {  
    case <-f.Released():  
        return nil  
    case <-ctx.Done():  
        return ctx.Err()  
    }  
}
```

optionally expose the channel

Waiting for a thing



```
var loaded Fence
var data map[string]int

err := Concurrently(ctx,
    func(ctx context.Context) error {
        defer loaded.Release()

        data, err := getData(ctx, url)
        if err != nil {
            return err
        }
        return nil
    },
```


Waiting for a thing



```
var loaded Fence
var data map[string]int

err := Concurrently(ctx,
    func(ctx context.Context) error {
        defer loaded.Release()


        data, err := getData(ctx, url)
        if err != nil {
            return err
        }
        return nil
    },
```



```
func(ctx context.Context) error {
    if err := loaded.Wait(ctx); err != nil {
        return err
    }
    return saveToCache(data)
},
func(ctx context.Context) error {
    if err := loaded.Wait(ctx); err != nil {
        return err
    }
    return processData(data)
})
```

Protecting State

Protecting State



```
func (s *ConcurrentState) With(
    ctx context.Context,
    fn func(*State) error,
) error {
    select {
    case state := <-s.state:
        defer func() { s.state <- state }()
        return fn(state)
    case ctx.Done():
        return ctx.Err()
    }
}
```

Protecting State



```
func (s *Listeners) AddListener(ctx context.Context, listener Listener) error {  
    return s.With(ctx, func(s *State) error {  
        s.listeners = append(s.listeners, listener)  
        return nil  
    })  
}
```

Limiters

Limiter



```
type Limiter struct {  
    limit chan struct{}  
    working sync.WaitGroup  
}  
  
func NewLimiter(n int) *Limiter {  
    return &Limiter{limit: make(chan struct{}), n}  
}
```

Limiter



```
type Limiter struct {  
    limit chan struct{}  
    working sync.WaitGroup  
}  
  
func NewLimiter(n int) *Limiter {  
    return &Limiter{limit: make(chan struct{}, n), working: sync.WaitGroup{}}
```



```
func (lim *Limiter) Go(ctx context.Context, fn func()) bool {  
    if ctx.Err() != nil {  
        return false  
    }  
  
    select {  
    case limiter.limit <- struct{}{}:  
    case <-ctx.Done():  
        return false  
    }  
  
    limiter.working.Add(1)  
    go func() {  
        defer func() {  
            <-limiter.limit  
            limiter.working.Done()  
        }()  
  
        fn()  
    }()  
  
    return true  
}  
  
func (lim *Limiter) Wait() {  
    limiter.working.Wait()  
}
```

Async processes in a server

Async processes in a server



```
func (server *Server) Run(ctx context.Context) error {
    server.pending = NewJobs(ctx)
    defer server.pending.Wait()

    // ... start listening on server ...
}

func (server *Server) ServeHTTP(w http.ResponseWriter, r *http.Request) {
    // ...
    started := server.pending.Go(r.Context(),
        func(ctx context.Context) {
            ... := server.db.ExecContext(ctx, ...)
        })
    if !started {
        // ... handle server shutting down or request canceled.
    }
    // ...
}
```

Async processes in a server

```
type Jobs struct {
    serverCtx context.Context
    group      errgroup.Group
}

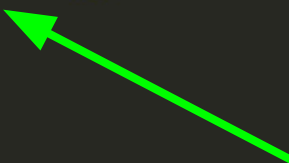
func NewJobs(serverCtx context.Context) *Jobs {
    return &Jobs{ serverCtx: serverCtx }
}

func (jobs *Jobs) Wait() { _ = jobs.group.Wait() }

func (jobs *Jobs) Go(requestCtx context.Context, fn func(ctx context.Context)) bool {
    if requestCtx.Err() != nil || jobs.serverCtx.Err() != nil {
        return false
    }

    jobs.group.Go(func() error {
        fn(jobs.serverCtx)
        return nil
    })

    return true
}
```



Async processes in a server (with limiter)



```
func (jobs *Jobs) Go(requestCtx context.Context, fn func(ctx context.Context)) bool {  
    if requestCtx.Err() != nil || jobs.serverCtx.Err() != nil {  
        return false  
    }  
  
    select {  
    case <-requestCtx.Done():  
        return false  
    case <-jobs.serverCtx.Done():  
        return false  
    case <-jobs.limiter:  
        defer func() { jobs.limiter <- struct{}{} }()  
    }  
  
    jobs.group.Go(func() error {  
        fn(jobs.serverCtx)  
        return nil  
    })  
  
    return true  
}
```


Additional resources:

- Bryan C. Mills - Rethinking Classical Concurrency Patterns
 - <https://www.youtube.com/watch?v=5zXAHh5tJqQ>
 - <https://drive.google.com/file/d/1nPdvhB0PutEJzdCq5ms6UI58dp50fcAN/view>
- Allen B. Downey - Little Book of Semaphores
 - <https://greenteapress.com/wp/semaphores/>
- Understanding Real-World Concurrency Bugs in Go
 - <https://songlh.github.io/paper/go-study.pdf>
- **#bestpractices** in Gophers Slack

Thank You